AWS Tutorial: AWS SageMaker

Train a Model with a Built-in Algorithm and Deploy It

Step 1: Create an Amazon SageMaker Notebook Instance

	Name	Instance	Creation time	Status	Actions
○	ExampleNotebookInstance	ml.t2.medium	Nov 13, 2018 03:47 UTC	InService	Open | Stop

Step 2: Create a Jupyter Notebook and Initialize Variables

```
In [ ]: from sagemaker import get_execution_role

        role = get_execution_role()
        bucket = 'test950607' # Use the name of your s3 bucket here
```

Step 3: Download, Explore, and Transform the Training Data

Download the MNIST Dataset

```
In [2]: %%time
        import pickle, gzip, numpy, urllib.request, json

        # Load the dataset
        urllib.request.urlretrieve("http://deeplearning.net/data/mnist/mnist.pkl.gz", "mnist.pkl.gz")
        with gzip.open('mnist.pkl.gz', 'rb') as f:
            train_set, valid_set, test_set = pickle.load(f, encoding='latin1')

CPU times: user 816 ms, sys: 321 ms, total: 1.14 s
Wall time: 2.66 s
```

Explore the Training Dataset

```
In [3]: %matplotlib inline
        import matplotlib.pyplot as plt
        plt.rcParams["figure.figsize"] = (2,10)


        def show_digit(img, caption='', subplot=None):
            if subplot == None:
                _, (subplot) = plt.subplots(1,1)
            imgr = img.reshape((28,28))
            subplot.axis('off')
            subplot.imshow(imgr, cmap='gray')
            plt.title(caption)

        show_digit(train_set[0][30], 'This is a {}'.format(train_set[1][30]))
```

Transform the Training Dataset and Upload It to S3

In [5]:
```
%%time
from sagemaker.amazon.common import write_numpy_to_dense_tensor
import io
import boto3

bucket = 'test950607' # Use the name of your s3 bucket here
data_key = 'kmeans_lowlevel_example/data'
data_location = 's3://{}/{}'.format(bucket, data_key)
print('training data will be uploaded to: {}'.format(data_location))

# Convert the training data into the format required by the SageMaker KMeans algorithm
buf = io.BytesIO()
write_numpy_to_dense_tensor(buf, train_set[0], train_set[1])
buf.seek(0)

boto3.resource('s3').Bucket(bucket).Object(data_key).upload_fileobj(buf)
```

```
training data will be uploaded to: s3://test950607/kmeans_lowlevel_example/data
CPU times: user 8.69 s, sys: 398 ms, total: 9.09 s
Wall time: 12 s
```

Step 4: Train a Model

Create a Training Job

In [6]:
```
from sagemaker import KMeans

data_location = 's3://{}/kmeans_highlevel_example/data'.format(bucket)
output_location = 's3://{}/kmeans_highlevel_example/output'.format(bucket)

print('training data will be uploaded to: {}'.format(data_location))
print('training artifacts will be uploaded to: {}'.format(output_location))

kmeans = KMeans(role=role,
                train_instance_count=2,
                train_instance_type='ml.c4.8xlarge',
                output_path=output_location,
                k=10,
                data_location=data_location)
```

```
training data will be uploaded to: s3://test950607/kmeans_highlevel_example/data
training artifacts will be uploaded to: s3://test950607/kmeans_highlevel_example/output
```

Start Train model

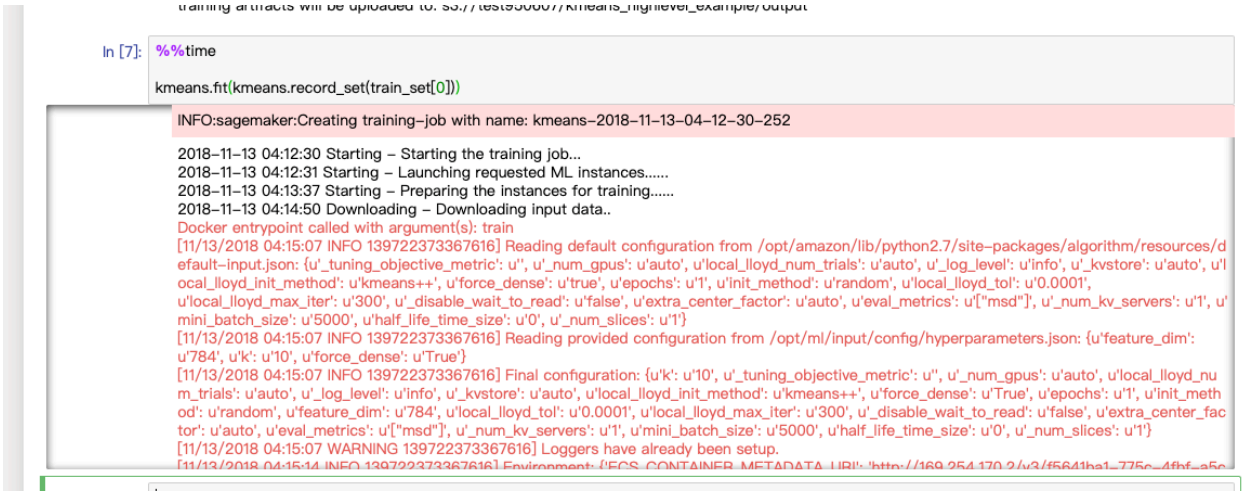

Use the SDK for Python

In [8]:

```
%%time
import boto3
from time import gmtime, strftime

job_name = 'kmeans-lowlevel-' + strftime("%Y-%m-%d-%H-%M-%S", gmtime())
print("Training job", job_name)

from sagemaker.amazon.amazon_estimator import get_image_uri
image = get_image_uri(boto3.Session().region_name, 'kmeans')

output_location = 's3://{}/kmeans_lowlevel_example/output'.format(bucket)
print('training artifacts will be uploaded to: {}'.format(output_location))

create_training_params = \
{
    "AlgorithmSpecification": {
        "TrainingImage": image,
        "TrainingInputMode": "File"
    },
    "RoleArn": role,
    "OutputDataConfig": {
        "S3OutputPath": output_location
    },
    "ResourceConfig": {
        "InstanceCount": 2,
        "InstanceType": "ml.c4.8xlarge",
        "VolumeSizeInGB": 50
    },
    "TrainingJobName": job_name,
    "HyperParameters": {
        "k": "10",
        "feature_dim": "784",
        "mini_batch_size": "500"
    },
    "StoppingCondition": {
        "MaxRuntimeInSeconds": 60 * 60
    },
    "InputDataConfig": [
        {
            "ChannelName": "train",
            "DataSource": {
                "S3DataSource": {
                    "S3DataType": "S3Prefix",
                    "S3Uri": data_location,
                    "S3DataDistributionType": "FullyReplicated"
                }
            },
```

Step 5: Deploy the Model to Amazon SageMaker Hosting Services

In [9]:
```
%%time

kmeans_predictor = kmeans.deploy(initial_instance_count=1,
                                 instance_type='ml.m4.xlarge')
```

```
INFO:sagemaker:Creating model with name: kmeans-2018-11-13-04-46-33-734
INFO:sagemaker:Creating endpoint with name kmeans-2018-11-13-04-12-30-252
```

```
--------------------------------------------------------------------!CPU times: user 305 ms, sys: 36.6 ms, total: 341 ms
Wall time: 5min 19s
```